

1.(b)

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
# Step 1: Generate a simple student marks dataset
```

```
np.random.seed(0)
```

```
marks = np.random.normal(loc=70, scale=10, size=30) # Normal distribution with mean=70, std=10
```

```
marks = np.clip(marks, 0, 100) # Ensure the marks are between 0 and 100
```

```
# Adding some random noise (outliers)
```

```
marks[5] = 120 # Outlier 1
```

```
marks[15] = -10 # Outlier 2
```

```
# Create a DataFrame
```

```
df = pd.DataFrame(marks, columns=['Marks'])
```

```
print("Original Marks Dataset with Noise:\n", df)
```

```
# Step 2: Outlier Detection and Removal using IQR method
```

```
Q1 = df['Marks'].quantile(0.25)
```

```
Q3 = df['Marks'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
# Remove outliers
```

```
df_cleaned = df[(df['Marks'] >= lower_bound) & (df['Marks'] <= upper_bound)]
```

```
# Step 3: Smoothing (Moving Average) to reduce minor noise
```

```
window_size = 5
```

```
df_cleaned['Smoothed Marks'] = df_cleaned['Marks'].rolling(window=window_size).mean()
```

```
# Plot the original and smoothed data
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(df_cleaned['Marks'].reset_index(drop=True), label="Cleaned Marks", marker='o')
```

```
plt.plot(df_cleaned['Smoothed Marks'].dropna().reset_index(drop=True), label="Smoothed Marks  
(Moving Avg)", marker='x')
```

```
plt.legend()
```

```
plt.title(f"Marks After Removing Outliers and Smoothing (Window Size: {window_size})")
```

```
plt.xlabel("Student Index")
```

```
plt.ylabel("Marks")
```

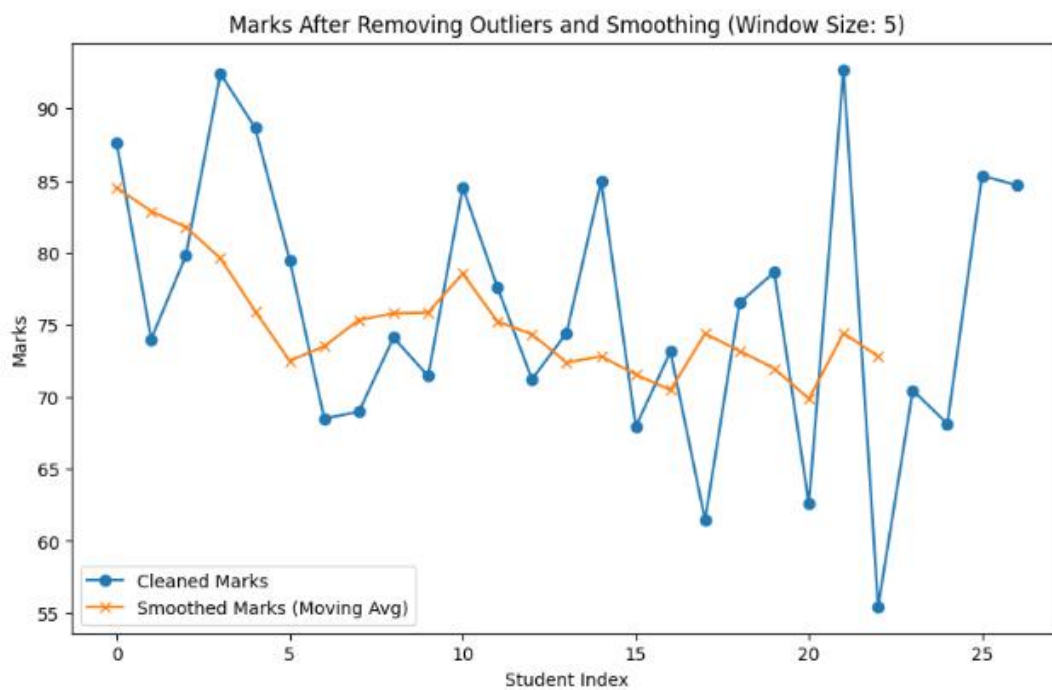
```
plt.show()
```

o\p

Original Marks Dataset with Noise:

	Marks
0	87.640523
1	74.001572
2	79.787380
3	92.408932
4	88.675580

5 120.000000
 6 79.500884
 7 68.486428
 8 68.967811
 9 74.105985
 10 71.440436
 11 84.542735
 12 77.610377
 13 71.216750
 14 74.438632
 15 -10.000000
 16 84.940791
 17 67.948417
 18 73.130677
 19 61.459043
 20 44.470102
 21 76.536186
 22 78.644362
 23 62.578350
 24 92.697546
 25 55.456343
 26 70.457585
 27 68.128161
 28 85.327792
 29 84.693588



2. Implement data processing to identify data redundancy and elimination

```
import pandas as pd

# Step 1: Create a simple student data dataset with some redundancy (duplicates)

data = {

    'StudentID': [101, 102, 103, 104, 105, 102, 106, 107, 105],

    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eva', 'Bob', 'Frank', 'Grace', 'Eva'],

    'Age': [20, 21, 22, 23, 24, 21, 25, 26, 24],

    'Grade': ['A', 'B', 'C', 'B', 'A', 'B', 'A', 'A', 'A']

}

# Create a DataFrame

df = pd.DataFrame(data)

# Display the original dataset with redundancy

print("Original Dataset with Redundancy (Duplicates):")

print(df)

# Step 2: Identify and eliminate redundant data (duplicate rows)

df_no_duplicates = df.drop_duplicates()

# Step 3: Display the cleaned dataset (duplicates removed)

print("\nCleaned Dataset (Duplicates Removed):")

print(df_no_duplicates)

#Step 4
```

```
Df_remove_praticular_duplicates_in_column=
```

```
df.drop_duplicates(subset='StudentID')
```

```
print(df_no_duplicates)
```

Original Dataset with Redundancy (Duplicates):

	StudentID	Name	Age	Grade
0	101	Alice	20	A
1	102	Bob	21	B
2	103	Charlie	22	C
3	104	David	23	B
4	105	Eva	24	A
5	102	Bob	21	B
6	106	Frank	25	A
7	107	Grace	26	A
8	105	Eva	24	A

Cleaned Dataset (Duplicates Removed):

	StudentID	Name	Age	Grade
0	101	Alice	20	A
1	102	Bob	21	B
2	103	Charlie	22	C
3	104	David	23	B
4	105	Eva	24	A
6	106	Frank	25	A
7	107	Grace	26	A

3. Implement any one imputation model

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.impute import SimpleImputer
```

```
# Step 1: Create a sample dataset with missing values (NaN)
```

```
data = {
```

```
    'StudentID': [101, 102, 103, 104, 105],
```

```
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eva'],
```

```
'Age': [20, np.nan, 22, 23, np.nan], # Missing values in 'Age'
'Grade': ['A', 'B', 'C', 'B', 'A']
}
```

```
# Create a DataFrame
```

```
df = pd.DataFrame(data)
```

```
# Display the original dataset with missing values
```

```
print("Original Dataset with Missing Values:")
```

```
print(df)
```

```
# Step 2: Impute missing values in the 'Age' column using SimpleImputer (mean strategy)
```

```
imputer = SimpleImputer(strategy='mean') # Use the 'mean' strategy for imputation
```

```
df['Age'] = imputer.fit_transform(df[['Age']]) # Impute missing values in the 'Age' column
```

```
# Display the dataset after imputation
```

```
print("\nDataset After Imputation (Mean):")
```

```
print(df)
```

```
output:
```

Original Dataset with Missing Values:

	StudentID	Name	Age	Grade
0	101	Alice	20.0	A
1	102	Bob	NaN	B
2	103	Charlie	22.0	C
3	104	David	23.0	B
4	105	Eva	NaN	A

Dataset After Imputation (Mean):

	StudentID	Name	Age	Grade
0	101	Alice	20.000000	A
1	102	Bob	21.666667	B
2	103	Charlie	22.000000	C
3	104	David	23.000000	B
4	105	Eva	21.666667	A

4.Implement Linear Regression

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Step 1: Create a simple dataset
```

```
# Let's create a simple dataset where we predict y based on x
```

```
data = {
```

```
    'X': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], # Input feature
```

```
    'y': [1.5, 3.2, 4.8, 6.5, 7.7, 9.1, 10.5, 11.6, 13.1, 14.3] # Target variable
```

```
}
```

```
# Convert the data into a pandas DataFrame
```

```
df = pd.DataFrame(data)
```

```
# Step 2: Visualize the data
```

```
plt.scatter(df['X'], df['y'], color='blue', label='Data points')
```

```
plt.xlabel('X')
```

```
plt.ylabel('y')
```

```
plt.title('Simple Linear Regression Example')
```

```
plt.show()
```

```
# Step 3: Split the data into training and testing sets
```

```
X = df[['X']] # Feature (independent variable)
```

```
y = df['y'] # Target (dependent variable)
```

```
# Split data into training (80%) and testing (20%)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 4: Initialize and train the Linear Regression model
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
# Step 5: Make predictions using the trained model
```

```
y_pred = model.predict(X_test)
```



```
# Step 6: Evaluate the model performance

mse = mean_squared_error(y_test, y_pred) # Mean Squared Error

r2 = r2_score(y_test, y_pred) # R-squared value


print(f'Mean Squared Error: {mse}')

print(f'R-squared: {r2}')


# Step 7: Visualize the regression line along with the data points

plt.scatter(X, y, color='blue', label='Data points') # Original data points

plt.plot(X, model.predict(X), color='red', label='Regression Line') # Regression line

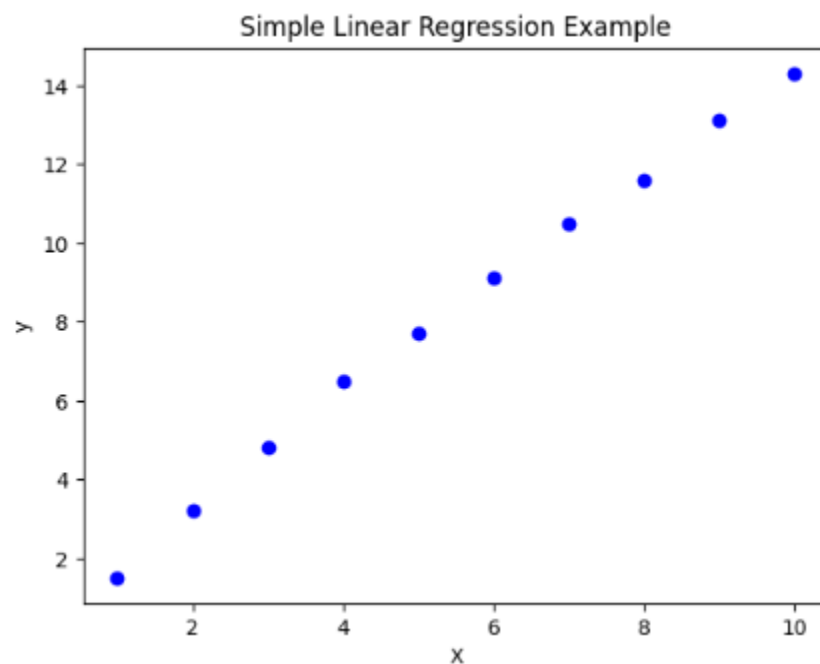
plt.xlabel('X')

plt.ylabel('Y')


plt.title('Linear Regression with Model Prediction')

plt.legend()

plt.show()
```



Mean Squared Error: 0.01193296670630199
R-squared: 0.999512989829352



Explanation:

1. **Dataset Creation:** We create a small dataset where x is the independent variable, and y is the dependent variable that we want to predict.
2. **Data Visualization:** We visualize the data with a scatter plot to understand the relationship between x and y .
3. **Train-Test Split:** We split the data into a training set (80%) and a testing set (20%) to evaluate the model's performance.
4. **Linear Regression Model:**
 - We use the `LinearRegression` class from `scikit-learn` to fit the model to the training data.
5. **Model Evaluation:** We evaluate the model's performance using:
 - **Mean Squared Error (MSE):** Measures the average squared difference between the predicted values and the actual values.
 - **R-squared:** Represents the proportion of variance in the dependent variable explained by the model.
6. **Regression Line:** Finally, we plot the regression line on top of the data points to visualize how well the model fits the data.

5. Implement logistic regression

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report


# Step 1: Create a simple dataset for binary classification

data = {

    'Study_Hours': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], # Number of hours studied

    'Passed': [0, 0, 0, 0, 1, 1, 1, 1, 1, 1] # 0=Fail, 1=Pass

}


# Convert the data into a pandas DataFrame

df = pd.DataFrame(data)


# Step 2: Visualize the data

plt.scatter(df['Study_Hours'], df['Passed'], color='blue', label='Data points')

plt.xlabel('Study Hours')

plt.ylabel('Passed (1) / Failed (0)')

plt.title('Logistic Regression Example')

plt.show()


# Step 3: Split the data into features (X) and target (y)
```

```
X = df[['Study_Hours']] # Feature (independent variable)

y = df['Passed']      # Target (dependent variable)


# Step 4: Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 5: Initialize and train the Logistic Regression model

model = LogisticRegression()

model.fit(X_train, y_train)


# Step 6: Make predictions using the trained model

y_pred = model.predict(X_test)


# Step 7: Evaluate the model performance

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

class_report = classification_report(y_test, y_pred)


# Output the evaluation results

print(f'Accuracy: {accuracy}')

print(f'Confusion Matrix:\n{conf_matrix}')

print(f'Classification Report:\n{class_report}')


# Step 8: Visualize the Logistic Regression decision boundary

# Plot the original data points
```

```
plt.scatter(df['Study_Hours'], df['Passed'], color='blue', label='Data points')
```

```
# Plot the logistic regression decision boundary
```

```
x_range = np.linspace(0, 10, 1000).reshape(-1, 1) # Create a range of study hours
```

```
y_range = model.predict_proba(x_range)[:, 1] # Get predicted probabilities for passing (class 1)
```

```
plt.plot(x_range, y_range, color='red', label='Decision Boundary (Probability)')
```


```
plt.xlabel('Study Hours')
```

```
plt.ylabel('Probability of Passing')
```

```
plt.title('Logistic Regression - Decision Boundary')
```

```
plt.legend()
```

```
plt.show()
```



Accuracy: 1.0

Confusion Matrix:


```
[[1 0]
```

```
 [0 1]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	1

accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2



6. Implement decision tree induction for classification

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier, plot_tree
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Step 1: Create a simple dataset for classification

```
data = {
```

```
    'Study_Hours': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], # Number of hours studied
```

```
    'Marks': [30, 35, 40, 50, 60, 65, 70, 80, 90, 95], # Marks obtained
```

```
    'Passed': [0, 0, 0, 0, 1, 1, 1, 1, 1, 1] # 0=Fail, 1=Pass
```

```
}
```

```
# Convert the data into a pandas DataFrame
```

```
df = pd.DataFrame(data)
```

```
# Step 2: Visualize the data
```

```
plt.scatter(df['Study_Hours'], df['Marks'], c=df['Passed'], cmap='coolwarm', marker='o')
```

```
plt.xlabel('Study Hours')
```

```
plt.ylabel('Marks')
```

```
plt.title('Decision Tree Classification Example')
```

```
plt.colorbar(label='Passed (1) / Failed (0)')
```

```
plt.show()
```

```
# Step 3: Split the data into features (X) and target (y)
```

```
X = df[['Study_Hours', 'Marks']] # Features (independent variables)
```

```
y = df['Passed'] # Target (dependent variable)
```

```
# Step 4: Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Step 5: Initialize and train the Decision Tree Classifier model
```

```
model = DecisionTreeClassifier(random_state=42)
```

```
model.fit(X_train, y_train)
```

```
# Step 6: Make predictions using the trained model
```

```
y_pred = model.predict(X_test)
```

```
# Step 7: Evaluate the model performance
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
class_report = classification_report(y_test, y_pred)
```

```
# Output the evaluation results
```

```
print(f'Accuracy: {accuracy}')
```

```
print(f'Confusion Matrix:\n{conf_matrix}')
```

```
print(f'Classification Report:\n{class_report}')
```


```
# Step 8: Visualize the Decision Tree
```

```
plt.figure(figsize=(12, 8))
```

```
plot_tree(model, filled=True, feature_names=['Study_Hours', 'Marks'], class_names=['Fail', 'Pass'],  
rounded=True, proportion=True)
```

```
plt.title('Decision Tree Classifier')
```

```
plt.show()
```

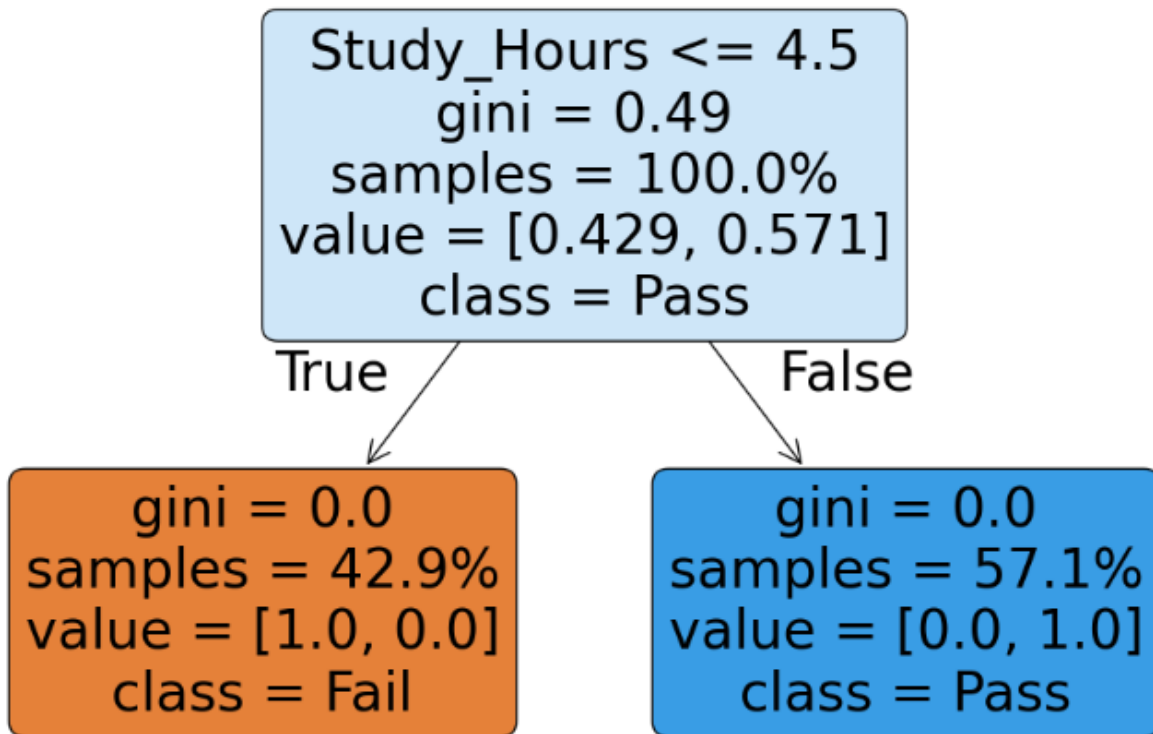



Accuracy: 1.0
Confusion Matrix:
[[1 0]
[0 2]]

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	2
accuracy			1.00	3
macro avg	1.00	1.00	1.00	3
weighted avg	1.00	1.00	1.00	3

Decision Tree Classifier



7. Implement random forest classifier

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Step 1: Create a simple dataset for classification

```
data = {
```

```
    'Study_Hours': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], # Number of hours studied
```

```
    'Marks': [30, 35, 40, 50, 60, 65, 70, 80, 90, 95], # Marks obtained
```

```
'Passed': [0, 0, 0, 0, 1, 1, 1, 1, 1, 1] # 0=Fail, 1=Pass
}

# Convert the data into a pandas DataFrame

df = pd.DataFrame(data)

# Step 2: Visualize the data

plt.scatter(df['Study_Hours'], df['Marks'], c=df['Passed'], cmap='coolwarm', marker='o')

plt.xlabel('Study Hours')

plt.ylabel('Marks')

plt.title('Random Forest Classification Example')

plt.colorbar(label='Passed (1) / Failed (0)')

plt.show()

# Step 3: Split the data into features (X) and target (y)

X = df[['Study_Hours', 'Marks']] # Features (independent variables)

y = df['Passed'] # Target (dependent variable)

# Step 4: Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Step 5: Initialize and train the Random Forest Classifier model

model = RandomForestClassifier(n_estimators=100, random_state=42)

model.fit(X_train, y_train)
```

Step 6: Make predictions using the trained model

```
y_pred = model.predict(X_test)
```

Step 7: Evaluate the model performance

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
class_report = classification_report(y_test, y_pred)
```

Output the evaluation results

```
print(f'Accuracy: {accuracy}')
```

```
print(f'Confusion Matrix:\n{conf_matrix}')
```

```
print(f'Classification Report:\n{class_report}')
```

Step 8: Visualize the Random Forest predictions

Plot the original data points

```
plt.scatter(df['Study_Hours'], df['Marks'], c=df['Passed'], cmap='coolwarm', marker='o')
```

Plot the predictions from the Random Forest model

Here, we will use the predicted probabilities of passing (class 1)

```
plt.scatter(X_test['Study_Hours'], X_test['Marks'], c=y_pred, marker='x', s=100, label='Predictions',  
edgecolor='black')
```

```
plt.xlabel('Study Hours')
```

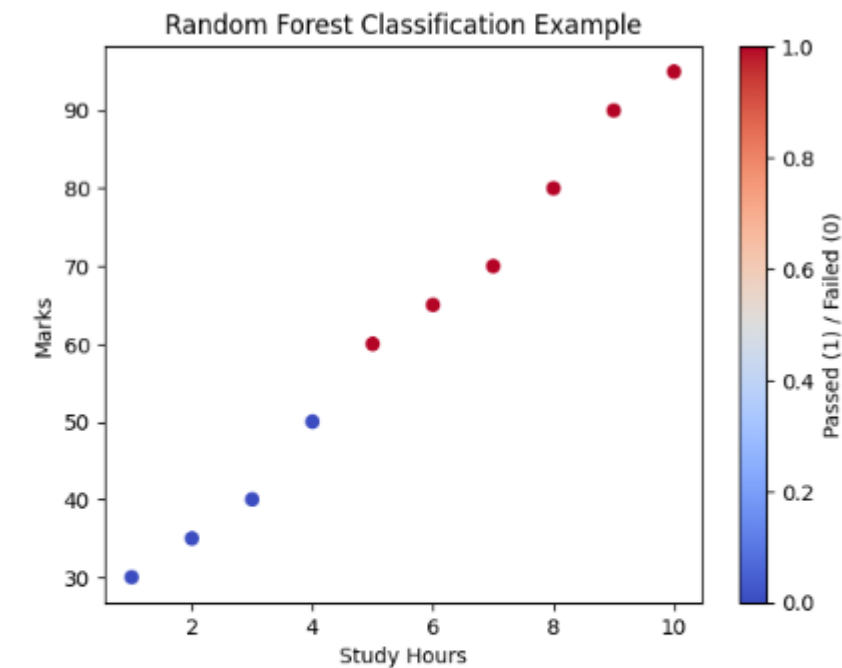
```
plt.ylabel('Marks')
```

```
plt.title('Random Forest Classifier - Predictions vs Actual')
```

```
plt.legend()
```

```
plt.colorbar(label='Passed (1) / Failed (0)')
```

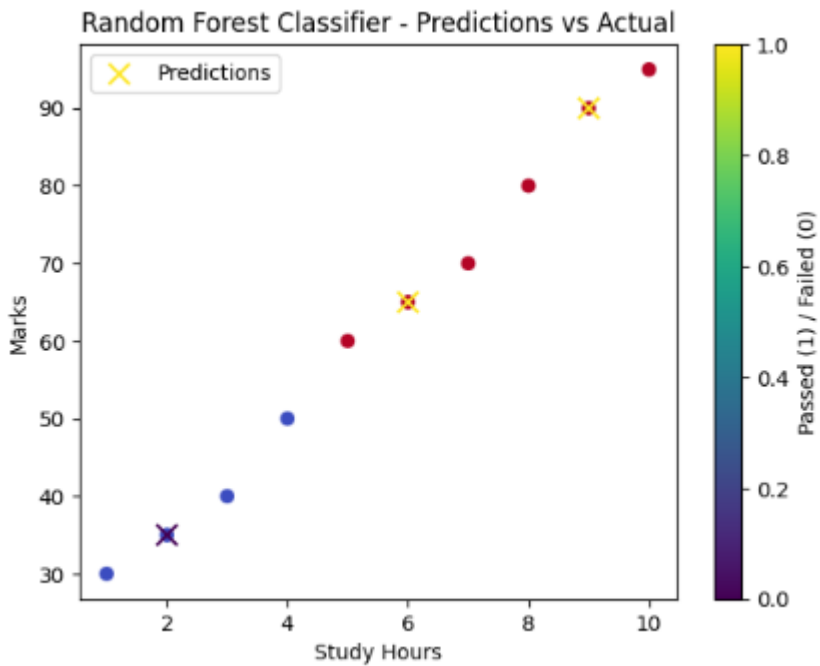
```
plt.show()
```



```
Accuracy: 1.0
Confusion Matrix:
[[1 0]
 [0 2]]
Classification Report:
              precision    recall  f1-score   support

     0       1.00      1.00      1.00         1
     1       1.00      1.00      1.00         2

 accuracy          1.00          1.00          1.00          3
 macro avg          1.00          1.00          1.00          3
weighted avg          1.00          1.00          1.00          3
```



8. Object segmentation using hierarchical based methods

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from skimage import io, color
```

```
from skimage.transform import resize
```

```
from scipy.cluster.hierarchy import linkage, fcluster
```

```
from skimage.segmentation import mark_boundaries
```

```
# Step 1: Load an Image
```

```
image = io.imread('FDP on Cybersecurity.jpg') # Example Image URL
```

```
image_rgb = image / 255.0 # Normalize image
```

```
# Downsample the image to reduce the size (e.g., resize to 1/4 of the original size)
```

```
downsampled_image = resize(image_rgb, (image_rgb.shape[0] // 4, image_rgb.shape[1] // 4),
mode='reflect')
```

```
# Show the original and downsampled image
```

```
plt.figure(figsize=(8, 6))
```

```
plt.subplot(1, 2, 1)
```

```
plt.imshow(image_rgb)
```

```
plt.title('Original Image')
```

```
plt.axis('off')
```

```
plt.subplot(1, 2, 2)
```

```
plt.imshow(downsampled_image)
```

```
plt.title('Downsampled Image')
```

```
plt.axis('off')
```

```
plt.show()
```

```
# Step 2: Pre-process Image (Reshape for clustering)
```

```
# Flatten the downsampled image
```

```
pixels = downsampled_image.reshape(-1, 3) # Shape: (number_of_pixels, 3)
```

```
# Step 3: Use only a subset of pixels for clustering
```

```
# Select a subset of pixels randomly (e.g., 10,000 pixels)
```

```
subset_size = 10000
```

```
np.random.seed(42) # For reproducibility
```

```
subset_indices = np.random.choice(pixels.shape[0], subset_size, replace=False)
```

```
subset_pixels = pixels[subset_indices]
```

```
# Step 4: Perform Hierarchical Clustering on the Subset of Pixels
```

```
Z = linkage(subset_pixels, method='ward') # 'ward' minimizes variance within clusters
```

```
# Step 5: Assign clusters (we define a threshold to segment the image)

num_clusters = 5 # You can change the number of clusters

clusters = fcluster(Z, num_clusters, criterion='maxclust')

# Now, we need to apply these clusters to the full downsampled image (not just the subset)

# Step 6: Map the clustering result back to the full image (since we're only using a subset, this step is
different)

# We will create an array that holds the labels for the entire image and set those corresponding to the
subset's indices.

# Create a full array of cluster labels for the image (it will have the same shape as the downsampled
image)

cluster_labels_full = np.zeros(pixels.shape[0], dtype=int)

# Assign the clusters to the labels corresponding to the subset indices

cluster_labels_full[subset_indices] = clusters

# Reshape the cluster labels to the shape of the downsampled image

segmented_image = cluster_labels_full.reshape(downsampled_image.shape[0],
downsampled_image.shape[1])

# Step 7: Visualize the Segmented Image

# Show the segmented image

plt.figure(figsize=(8, 6))

plt.imshow(segmented_image, cmap='jet') # Use 'jet' color map for better visualization

plt.title('Hierarchical Segmentation (Clustering)')

plt.axis('off')

plt.show()

# Step 8: Mark boundaries on the downsampled image

# Use mode='thick' to mark the boundaries
```



```
boundaries = mark_boundaries(downsampled_image, segmented_image, color=(1, 0, 0), mode='thick')
```

```
# Show image with boundaries marked
```

```
plt.figure(figsize=(8, 6))
```


```
plt.imshow(boundaries)
```

```
plt.title('Boundaries of Segments')
```


```
plt.axis('off')
```

```
plt.show()
```


OUTPUT:



Hierarchical Segmentation (Clustering)



Boundaries of Segments



9. perform visualization techniques like bar,column ,line,scatter,3d cubes

```
import matplotlib.pyplot as plt

import seaborn as sns

import numpy as np

from mpl_toolkits.mplot3d import Axes3D


# Sample data for the visualizations

categories = ['A', 'B', 'C', 'D', 'E']

values = [10, 20, 15, 25, 30]

x = np.linspace(0, 10, 100) # For line and scatter plots

y = np.sin(x) # Line plot data

z = np.cos(x) # For scatter plot

z3d = np.random.rand(5) # 3D cube plot data


# ----- Bar Chart -----

plt.figure(figsize=(8, 6))

plt.bar(categories, values, color='skyblue')

plt.xlabel('Categories')

plt.ylabel('Values')

plt.title('Bar Chart')

plt.show()
```

----- Column Chart -----

```
plt.figure(figsize=(8, 6))  
  
plt.barh(categories, values, color='lightcoral')  
  
plt.xlabel('Values')  
  
plt.ylabel('Categories')  
  
plt.title('Column Chart (Horizontal Bar)')  
  
plt.show()
```

----- Line Chart -----

```
plt.figure(figsize=(8, 6))  
  
plt.plot(x, y, label='sin(x)', color='blue')  
  
plt.plot(x, z, label='cos(x)', color='green')  
  
plt.xlabel('X-axis')  
  
plt.ylabel('Y-axis')  
  
plt.title('Line Chart')  
  
plt.legend()  
  
plt.show()
```

----- Scatter Plot -----


```
plt.figure(figsize=(8, 6))

plt.scatter(x, y, color='red', label='sin(x)')

plt.scatter(x, z, color='purple', label='cos(x)')

plt.xlabel('X-axis')

plt.ylabel('Y-axis')

plt.title('Scatter Plot')

plt.legend()

plt.show()
```

```
# ----- 3D Cube Plot (3D Scatter) -----
```

```
fig = plt.figure(figsize=(8, 6))

ax = fig.add_subplot(111, projection='3d')

# Data for 3D plot (random cubes)

x3d = np.random.rand(5)

y3d = np.random.rand(5)

z3d = np.random.rand(5)

# 3D scatter plot

ax.scatter(x3d, y3d, z3d, c='r', marker='o')

# Label axes

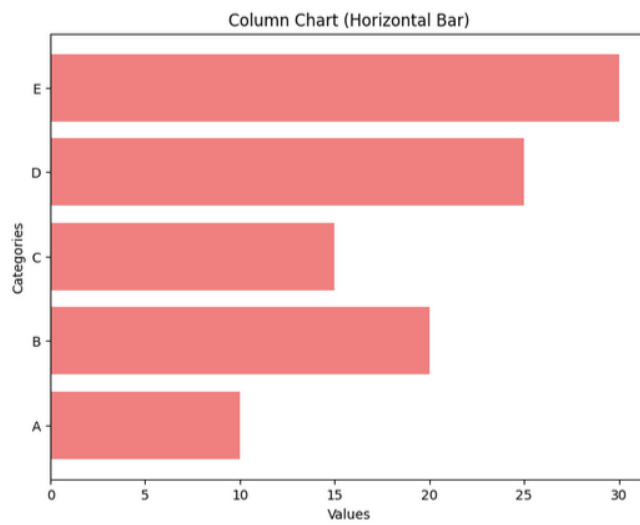

ax.set_xlabel('X Label')
```




```
ax.set_ylabel('Y Label')
```

```
ax.set_zlabel('Z Label')
```

```
ax.set_title('3D Cube (Scatter Plot)')
```

```
plt.show()
```





3D Cube (Scatter Plot)

10.Perform Descriptive analysis on health care data .

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns


# Create a sample healthcare dataset

data = {

    'Age': np.random.randint(20, 80, 100), # Age between 20 and 80

    'Gender': np.random.choice(['Male', 'Female'], 100),

    'Blood Pressure': np.random.randint(90, 180, 100), # Blood Pressure between
90 and 180

    'Cholesterol': np.random.randint(100, 300, 100), # Cholesterol between 100
and 300

    'Heart Disease': np.random.choice([0, 1], 100), # 0 = No, 1 = Yes

}


# Create a DataFrame

df = pd.DataFrame(data)


# Descriptive statistics for numerical columns

print(df.describe())
```

```
# Visualize the distribution of Age and Blood Pressure
```

```
plt.figure(figsize=(10, 5))
```

```
sns.histplot(df['Age'], kde=True, color='blue', label='Age')
```

```
sns.histplot(df['Blood Pressure'], kde=True, color='green', label='Blood Pressure')
```

```
plt.legend()
```

```
plt.title('Age and Blood Pressure Distribution')
```

```
plt.show()
```

```
# Countplot for Gender distribution
```

```
plt.figure(figsize=(6, 4))
```

```
sns.countplot(x='Gender', data=df, palette='Set1')
```

```
plt.title('Gender Distribution')
```

```
plt.show()
```

```
# Correlation heatmap for numerical variables only (Age, Blood Pressure,  
Cholesterol)
```

```
numeric_columns = df.select_dtypes(include=[np.number]) # Select only numeric  
columns
```



```
plt.figure(figsize=(8, 6))
```

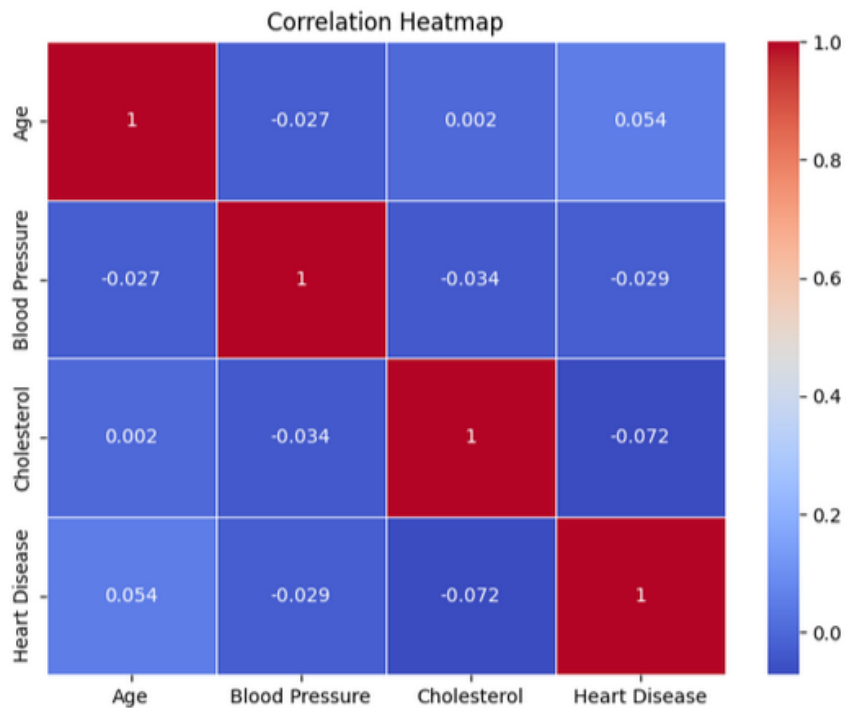
```
sns.heatmap(numeric_columns.corr(), annot=True, cmap='coolwarm',  
linewidths=0.5)
```

```
plt.title('Correlation Heatmap')
```

```
plt.show()
```

	Age	Blood Pressure	Cholesterol	Heart Disease
count	100.00000	100.00000	100.00000	100.00000
mean	51.44000	135.74000	216.27000	0.47000
std	17.00821	25.22689	61.35916	0.50161
min	20.00000	90.00000	100.00000	0.00000
25%	35.75000	113.00000	157.75000	0.00000
50%	52.50000	138.00000	227.00000	0.00000
75%	66.00000	156.25000	273.75000	1.00000
max	79.00000	179.00000	299.00000	1.00000





11. Perform Predictive analytics on Product Sales data

Step 1: Import Necessary Libraries

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

Step 2: Create a Mock Product Sales Dataset

```
data = {
```

```
&#39;Price&#39;: [15, 20, 25, 30, 35, 40, 45, 50, 55, 60],  
  
&#39;Advertising&#39;: [1000, 2000, 3000, 4000, 5000, 6000, 7000, 8000, 9000,  
10000],  
  
&#39;Season&#39;: [&#39;Spring&#39;, &#39;Summer&#39;, &#39;Fall&#39;,,  
&#39;Winter&#39;, &#39;Spring&#39;, &#39;Summer&#39;, &#39;Fall&#39;,,  
&#39;Winter&#39;, &#39;Spring&#39;, &#39;Summer&#39;],  
  
&#39;Sales&#39;: [150, 180, 160, 140, 130, 120, 110, 115, 200, 210]  
  
}
```

```
# Convert the dictionary into a pandas DataFrame
```

```
df = pd.DataFrame(data)
```

```
# Step 3: Preprocess the Data
```

```
# Convert &#39;Season&#39; to numeric values (Spring = 0, Summer = 1, Fall = 2,  
Winter = 3)
```

```
df[&#39;Season&#39;] = df[&#39;Season&#39;].map({&#39;Spring&#39;: 0,  
&#39;Summer&#39;: 1, &#39;Fall&#39;: 2, &#39;Winter&#39;: 3})
```

```
# Step 4: Select Features and Target Variable
```

```
X = df[ [&#39;Price&#39;, &#39;Advertising&#39;, &#39;Season&#39;]] # Features
```

```
y = df[&#39;Sales&#39;] # Target variable (Sales)
```

Step 5: Split the Data into Training and Testing Sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

Step 6: Train the Linear Regression Model

```
model = LinearRegression()  
  
model.fit(X_train, y_train)
```

Step 7: Make Predictions

```
y_pred = model.predict(X_test)
```

Step 8: Evaluate the Model

```
mse = mean_squared_error(y_test, y_pred)  
  
r2 = r2_score(y_test, y_pred)
```

Print Evaluation Metrics

```
print(f'&#39;Mean Squared Error: {mse}&#39;)  
  
print(f'&#39;R-squared: {r2}&#39;)
```

Step 9: Visualize the Results

```
# Plot Actual vs Predicted Sales
```

```
plt.scatter(y_test, y_pred)
```

```
plt.xlabel('#39;Actual Sales#39;)
```

```
plt.ylabel('#39;Predicted Sales#39;)
```

```
plt.title('#39;Actual vs Predicted Sales#39;)
```

```
plt.show()
```

```
# Step 10: Model Coefficients
```

```
print('&quot;\nModel Coefficients:&quot;)
```


```
coefficients = pd.DataFrame(model.coef_, X.columns,  
columns=['#39;Coefficient#39;])
```

```
print(coefficients)
```

OUTPUT:

Mean Squared Error: 1593.2661115916937

R-squared: -14.932661115916938



Model Coefficients:

	Coefficient
Price	0.000014
Advertising	0.002721
Season	-8.382353

12. Apply Predictive analytics for Weather forecasting.

Step 1: Import Necessary Libraries

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

Step 2: Create a Mock Weather Dataset

```
data = {  
    'Temperature': [30, 32, 33, 31, 29, 28, 25, 27, 30, 31, 33, 35, 36, 37, 34],  
    'Humidity': [80, 75, 77, 70, 85, 88, 90, 85, 80, 78, 76, 74, 73, 72, 71],  
    'Wind Speed': [10, 12, 15, 11, 13, 14, 9, 10, 12, 11, 10, 9, 8, 7, 6],  
    'Pressure': [1010, 1012, 1011, 1010, 1011, 1013, 1012, 1010, 1011, 1012,  
1013, 1014, 1015, 1016, 1017],  
    'Month': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 1, 2, 3]  
}
```

Convert the dictionary into a pandas DataFrame

```
df = pd.DataFrame(data)
```

```
# Step 3: Explore the Dataset
```

```
print(df.head())
```

```
print("\nSummary Statistics:")
```

```
print(df.describe())
```

```
# Step 4: Handle Missing Values (not needed in this mock dataset)
```

```
# df.fillna(df.median(), inplace=True)
```

```
# Step 5: Select Features and Target Variable
```

```
X = df[['Humidity', 'Wind Speed', 'Pressure', 'Month']] # Features
```

```
y = df['Temperature'] # Target variable (Temperature)
```

```
# Step 6: Split the Data into Training and Testing Sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

```
# Step 7: Train the Linear Regression Model
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

Step 8: Make Predictions

```
y_pred = model.predict(X_test)
```

Step 9: Evaluate the Model

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f"\nMean Squared Error: {mse}")
```

```
print(f"R-squared: {r2}")
```

Step 10: Visualize the Results

Plot Actual vs Predicted Temperature

```
plt.scatter(y_test, y_pred)
```

```
plt.xlabel('Actual Temperature')
```

```
plt.ylabel('Predicted Temperature')
```

```
plt.title('Actual vs Predicted Temperature')
```

```
plt.show()
```

Step 11: Model Interpretation

```
# Display model coefficients
```

```
coefficients = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
```

```
print("\nModel Coefficients:")
```

```
print(coefficients)
```


	Temperature	Humidity	Wind Speed	Pressure	Month
0	30	80	10	1010	1
1	32	75	12	1012	2
2	33	77	15	1011	3
3	31	70	11	1010	4
4	29	85	13	1011	5

```
Summary Statistics:
```

	Temperature	Humidity	Wind Speed	Pressure	Month
count	15.000000	15.000000	15.000000	15.000000	15.000000
mean	31.400000	78.266667	10.466667	1012.466667	5.600000
std	3.376389	6.284524	2.503331	2.199567	3.718679
min	25.000000	70.000000	6.000000	1010.000000	1.000000
25%	29.500000	73.500000	9.000000	1011.000000	2.500000
50%	31.000000	77.000000	10.000000	1012.000000	5.000000
75%	33.500000	82.500000	12.000000	1013.500000	8.500000
max	37.000000	90.000000	15.000000	1017.000000	12.000000

```
Mean Squared Error: 1.6787344343422512
```

```
R-squared: 0.6402711926409461
```



Model Coefficients:

	Coefficient
Humidity	-0.388341
Wind Speed	0.377432
Pressure	0.837989
Month	-0.013751